

10-701: Introduction to Machine Learning

Lecture 7 –Logistic Regression

Hoda Heidari

* Slides adopted from F24 offering of 10701 by Henry Chai.

Recall: Probabilistic Learning

- Previously:
 - (Unknown) Target function, $c^*: \mathcal{X} \rightarrow \mathcal{Y}$
 - Classifier, $h: \mathcal{X} \rightarrow \mathcal{Y}$
 - Goal: find a classifier, h , that best approximates c^*
- Now:
 - (Unknown) Target *distribution*, $y \sim P^*(Y|\mathbf{x})$
 - Distribution, $P(Y|\mathbf{x})$
 - Goal: find a distribution, P , that best approximates P^*

Recipe for Naïve Bayes

- Define a model space and model parameters
 - Make the Naïve Bayes assumption

$$P(X|Y) = \prod_{d=1}^D P(X_d|Y)$$

- Parameters: $\pi = P(Y = 1)$, $\theta_{d,y} = P(X_d = 1|Y = y)$
- Write down an objective function
 - Maximize the log-likelihood
- Optimize the objective w.r.t. the model parameters
 - Solve in *closed form*: take partial derivatives, set to 0 and solve

Bernoulli Naïve Bayes

- Binary label
 - $Y \sim \text{Bernoulli}(\pi)$
 - $\hat{\pi} = N_{Y=1} / N$
 - N = # of data points
 - $N_{Y=1}$ = # of data points with label 1
- Binary features
 - $X_d | Y = y \sim \text{Bernoulli}(\theta_{d,y})$
 - $\hat{\theta}_{d,y} = N_{Y=y, X_d=1} / N_{Y=y}$
 - $N_{Y=y}$ = # of data points with label y
 - $N_{Y=y, X_d=1}$ = # of data points with label y and feature $X_d = 1$

Bernoulli Naïve Bayes: Making Predictions

- Given a test data point $\mathbf{x}' = [x'_1, \dots, x'_D]^T$

Bernoulli Naïve Bayes: Making Predictions

- Given a test data point $\mathbf{x}' = [x'_1, \dots, x'_D]^T$

$$P(Y = 1|\mathbf{x}') \propto P(Y = 1)P(\mathbf{x}'|Y = 1)$$

$$= \hat{\pi} \prod_{d=1}^D \hat{\theta}_{d,1}^{x'_d} (1 - \hat{\theta}_{d,1})^{1-x'_d}$$

$$P(Y = 0|\mathbf{x}') \propto (1 - \hat{\pi}) \prod_{d=1}^D \hat{\theta}_{d,0}^{x'_d} (1 - \hat{\theta}_{d,0})^{1-x'_d}$$

$$\hat{y} = \begin{cases} 1 & \text{if } \hat{\pi} \prod_{d=1}^D \hat{\theta}_{d,1}^{x'_d} (1 - \hat{\theta}_{d,1})^{1-x'_d} > \\ & (1 - \hat{\pi}) \prod_{d=1}^D \hat{\theta}_{d,0}^{x'_d} (1 - \hat{\theta}_{d,0})^{1-x'_d} \\ 0 & \text{otherwise} \end{cases}$$

What if some
Word-Label
pair never
appears in our
training data?

x_1 ("hat")	x_2 ("cat")	x_3 ("dog")	x_4 ("fish")	x_5 ("mom")	x_6 ("dad")	y (Dr. Seuss)
1	1	0	0	0	0	1
0	0	1	0	0	0	0
0	0	0	1	0	0	1
0	0	0	0	1	0	0

The Cat in the Hat gets a Dog (by ???)

- If some $\hat{\theta}_{d,y} = 0$ and that word appears in our test data \mathbf{x}' , then $P(Y = y|\mathbf{x}') = 0$ even if all the other features in \mathbf{x}' point to the label being y !
- The model has been overfit to the training data...
- We can address this with a prior over the parameters!

Setting the Parameters via MAP

- Binary label
 - $Y \sim \text{Bernoulli}(\pi)$
 - $\hat{\pi} = N_{Y=1} / N$
 - $N = \#$ of data points
 - $N_{Y=1} = \#$ of data points with label 1
- Binary features
 - $X_d | Y = y \sim \text{Bernoulli}(\theta_{d,y})$ and $\theta_{d,y} \sim \text{Beta}(\alpha, \beta)$
 - $\hat{\theta}_{d,y} = N_{Y=y, X_d=1} + (\alpha - 1) / N_{Y=y} + (\alpha - 1) + (\beta - 1)$
 - $N_{Y=y} = \#$ of data points with label y
 - $N_{Y=y, X_d=1} = \#$ of data points with label y and feature $X_d = 1$
 - α and β are “pseudocounts” of imagined data points that help avoid zero-probability predictions.
 - Common choice: $\alpha = \beta = 2$

What can we do when this is a bad/incorrect assumption, e.g., when our features are words in a sentence?

- **Assume** features are conditionally independent given the label:

$$P(X|Y) = \prod_{d=1}^D P(X_d|Y)$$

- Pros:
 - Significantly reduces computational complexity
 - Also reduces model complexity, combats overfitting
- Cons:
 - Is a strong, often illogical assumption
 - We'll see a relaxed version of this much later when we discuss Bayesian networks

Key Takeaways

- Text data
 - Bag-of-words feature representation
- Naïve Bayes
 - Conditional independence assumption
 - Pros and cons
 - Different Naïve Bayes models based on type of features
 - MLE vs. MAP for Bernoulli Naïve Bayes

Recall: Building a Probabilistic Classifier

- Define a decision rule
 - Given a test data point \mathbf{x}' , predict its label \hat{y} using the *posterior distribution* $P(Y = y|X = \mathbf{x}')$
 - Common choice: $\hat{y} = \underset{y}{\operatorname{argmax}} P(Y = y|X = \mathbf{x}')$
- Model the posterior distribution
 - Option 1 - Model $P(Y|X)$ directly as some function of X (today!)
 - Option 2 - Use Bayes' rule (Monday):

$$P(Y|X) = \frac{P(X|Y) P(Y)}{P(X)} \propto P(X|Y) P(Y)$$

Modelling the Posterior

- Suppose we have binary labels $y \in \{0,1\}$ and D -dimensional inputs $\mathbf{x} = [1, x_1, \dots, x_D]^T \in \mathbb{R}^{D+1}$
- **Assume**
- This implies two useful facts:

Modelling the Posterior

- Suppose we have binary labels $y \in \{0,1\}$ and D -dimensional inputs $\mathbf{x} = [1, x_1, \dots, x_D]^T \in \mathbb{R}^{D+1}$


- **Assume**

$$\begin{aligned} P(Y = 1|\mathbf{x}) &= \text{logistic}(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x})} \\ &= \frac{\exp(\mathbf{w}^T \mathbf{x})}{\exp(\mathbf{w}^T \mathbf{x}) + 1} \end{aligned}$$


- This implies two useful facts:

1. $P(Y = 0|\mathbf{x}) = 1 - P(Y = 1|\mathbf{x}) = \frac{1}{\exp(\mathbf{w}^T \mathbf{x}) + 1}$
2. $\frac{P(Y = 1|\mathbf{x})}{P(Y = 0|\mathbf{x})} = \exp(\mathbf{w}^T \mathbf{x}) \rightarrow \log \frac{P(Y = 1|\mathbf{x})}{P(Y = 0|\mathbf{x})} = \mathbf{w}^T \mathbf{x}$

Logistic Function



Why use the Logistic Function?



- Differentiable everywhere
- $\text{logistic}: \mathbb{R} \rightarrow [0, 1]$

Logistic Regression Decision Boundary

The decision boundary is linear in x !

Logistic Regression Decision Boundary

The decision boundary is linear in \mathbf{x} !

$$\hat{y} = \begin{cases} 1 & \text{if } P(Y = 1|\mathbf{x}) \geq \frac{1}{2} \\ 0 & \text{otherwise} \end{cases}$$

$$P(Y = 1|\mathbf{x}) = \text{logistic}(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x})} \geq \frac{1}{2}$$


$$2 \geq 1 + \exp(-\mathbf{w}^T \mathbf{x})$$

$$1 \geq \exp(-\mathbf{w}^T \mathbf{x})$$


$$\log(1) \geq -\mathbf{w}^T \mathbf{x}$$

$$0 \leq \mathbf{w}^T \mathbf{x}$$


Logistic Regression Decision Boundary



Logistic Regression Decision Boundary



Logistic Regression Decision Boundary



General Recipe for Machine Learning

- Define a model space and model parameters
- Write down an objective function
- Optimize the objective w.r.t. the model parameters

Recipe for Logistic Regression

- Define a model space and model parameters
 - Assume independent, identically distributed (iid) data
 - Assume $P(Y = 1|X) = \text{logistic}(\mathbf{w}^T \mathbf{x})$
 - Parameters: $\mathbf{w} = [w_0, w_1, \dots, w_D]$
- Write down an objective function
 - ~~Maximize the conditional log-likelihood~~
 - Minimize the negative conditional log-likelihood
- Optimize the objective w.r.t. the model parameters
 - ???

Setting the Parameters via Minimum Negative Conditional (log-)Likelihood Estimation (MCLE)

Find \mathbf{w} that minimizes

$$\ell_{\mathcal{D}}(\mathbf{w}) =$$

Setting the Parameters via Minimum Negative Conditional (log-)Likelihood Estimation (MCLE)

Find \mathbf{w} that minimizes

$$\begin{aligned}\ell_{\mathcal{D}}(\mathbf{w}) &= -\log P(y^{(1)}, \dots, y^{(N)} | \mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}, \mathbf{w}) = -\log \prod_{n=1}^N P(y^{(n)} | \mathbf{x}^{(n)}, \mathbf{w}) \\ &= -\log \prod_{n=1}^N P(Y = 1 | \mathbf{x}^{(n)}, \mathbf{w})^{y^{(n)}} \left(P(Y = 0 | \mathbf{x}^{(n)}, \mathbf{w}) \right)^{1-y^{(n)}} \\ &= -\sum_{n=1}^N y^{(n)} \log P(Y = 1 | \mathbf{x}^{(n)}, \mathbf{w}) + (1 - y^{(n)}) \log P(Y = 0 | \mathbf{x}^{(n)}, \mathbf{w}) \\ &= -\sum_{n=1}^N y^{(n)} \log \frac{P(Y = 1 | \mathbf{x}^{(n)}, \mathbf{w})}{P(Y = 0 | \mathbf{x}^{(n)}, \mathbf{w})} + \log P(Y = 0 | \mathbf{x}^{(n)}, \mathbf{w}) \\ &= -\sum_{n=1}^N y^{(n)} \mathbf{w}^T \mathbf{x}^{(n)} - \log \left(1 + \exp(\mathbf{w}^T \mathbf{x}^{(n)}) \right)\end{aligned}$$

Minimizing the Negative Conditional (log-)Likelihood


$$\ell_{\mathcal{D}}(\mathbf{w}) = - \sum_{n=1}^N y^{(n)} \mathbf{w}^T \mathbf{x}^{(n)} - \log \left(1 + \exp(\mathbf{w}^T \mathbf{x}^{(n)}) \right)$$

Minimizing the Negative Conditional (log-)Likelihood

$$\begin{aligned}\ell_{\mathcal{D}}(\mathbf{w}) &= - \sum_{n=1}^N y^{(n)} \mathbf{w}^T \mathbf{x}^{(n)} - \log \left(1 + \exp(\mathbf{w}^T \mathbf{x}^{(n)}) \right) \\ \nabla_{\mathbf{w}} \ell_{\mathcal{D}}(\mathbf{w}) &= - \sum_{n=1}^N y^{(n)} \nabla_{\mathbf{w}} \mathbf{w}^T \mathbf{x}^{(n)} - \nabla_{\mathbf{w}} \log \left(1 + \exp(\mathbf{w}^T \mathbf{x}^{(n)}) \right) \\ &= - \sum_{n=1}^N y^{(n)} \mathbf{x}^{(n)} - \frac{\exp(\mathbf{w}^T \mathbf{x}^{(n)})}{1 + \exp(\mathbf{w}^T \mathbf{x}^{(n)})} \mathbf{x}^{(n)} \\ &= \sum_{n=1}^N \mathbf{x}^{(n)} (P(Y = 1 | \mathbf{x}^{(n)}, \mathbf{w}) - y^{(n)})\end{aligned}$$


Recall: Gradient Descent

- An iterative method for minimizing functions
- Requires the gradient to exist everywhere



Recall: Gradient Descent

- An iterative method for minimizing functions
- Requires the gradient to exist everywhere



- Good news: the negative conditional log-likelihood, like the squared error, is also convex!

Gradient Descent

- Input: $\mathcal{D} = \{(\mathbf{x}^{(n)}, y^{(n)})\}_{n=1}^N, \eta^{(0)}$

1. Initialize $\mathbf{w}^{(0)}$ to all zeros and set $t = 0$

2. While TERMINATION CRITERION is not satisfied

- a. Compute the gradient:

$$O(ND) \left\{ \nabla_{\mathbf{w}} \ell_{\mathcal{D}}(\mathbf{w}^{(t)}) = \sum_{n=1}^N \mathbf{x}^{(n)} (P(Y = 1 | \mathbf{x}^{(n)}, \mathbf{w}^{(t)}) - y^{(n)}) \right.$$

- b. Update \mathbf{w} : $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \eta^{(0)} \nabla_{\mathbf{w}} \ell_{\mathcal{D}}(\mathbf{w}^{(t)})$

- c. Increment t : $t \leftarrow t + 1$

- Output: $\mathbf{w}^{(t)}$

Stochastic Gradient Descent

- Input: $\mathcal{D} = \{(\mathbf{x}^{(n)}, y^{(n)})\}_{n=1}^N, \eta_{SGD}^{(0)}$
 1. Initialize $\mathbf{w}^{(0)}$ to all zeros and set $t = 0$
 2. While TERMINATION CRITERION is not satisfied
 - a. Randomly sample a data point from $\mathcal{D}, (\mathbf{x}^{(n)}, y^{(n)})$
 - b. Compute the pointwise gradient:
$$\nabla_{\mathbf{w}} \ell^{(n)}(\mathbf{w}^{(t)}) = \mathbf{x}^{(n)} (P(Y = 1 | \mathbf{x}^{(n)}, \mathbf{w}^{(t)}) - y^{(n)})$$
 - c. Update \mathbf{w} : $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \eta_{SGD}^{(0)} \nabla_{\mathbf{w}} \ell^{(n)}(\mathbf{w}^{(t)})$
 - d. Increment t : $t \leftarrow t + 1$
- Output: $\mathbf{w}^{(t)}$


Stochastic Gradient Descent

- If the data point is sampled uniformly at random, then the expected value of the pointwise gradient is proportional to the full gradient:


$$\begin{aligned} E \left[\nabla_{\mathbf{w}} \ell_{\mathbf{x}^{(n)}, \mathbf{y}^{(n)}}(\mathbf{w}^{(t)}) \right] &= \frac{1}{N} \sum_{n=1}^N \nabla_{\mathbf{w}} \ell^{(n)}(\mathbf{w}^{(t)}) \\ &= \frac{1}{N} \sum_{n=1}^N \mathbf{x}^{(n)} (P(Y = 1 | \mathbf{x}^{(n)}, \mathbf{w}^{(t)}) - y^{(n)}) \\ &= \frac{1}{N} \nabla_{\mathbf{w}} \ell_{\mathcal{D}}(\mathbf{w}^{(t)}) \end{aligned}$$

- In practice, the data set is randomly shuffled then looped through so that each data point is used equally often

Stochastic Gradient Descent vs. Gradient Descent



Gradient Descent



Stochastic Gradient Descent

Mini-batch Stochastic Gradient Descent

- Input: $\mathcal{D} = \{(\mathbf{x}^{(n)}, y^{(n)})\}_{n=1}^N, \eta_{MB}^{(0)}, B$

1. Initialize $\mathbf{w}^{(0)}$ to all zeros and set $t = 0$
2. While TERMINATION CRITERION is not satisfied

- a. Randomly sample B data points from \mathcal{D} :

$$\mathcal{D}_{batch} \{(\mathbf{x}^{(b)}, y^{(b)})\}_{b=1}^B$$

- b. Compute the gradient w.r.t. the sampled *batch*:


$$\nabla_{\mathbf{w}} \ell_{\mathcal{D}_{batch}}(\mathbf{w}^{(t)}) = \sum_{b=1}^B \mathbf{x}^{(b)} (P(Y = 1 | \mathbf{x}^{(b)}, \mathbf{w}) - y^{(b)})$$

- c. Update \mathbf{w} : $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \eta_{MB}^{(0)} \nabla_{\mathbf{w}} \ell_{\mathcal{D}_{batch}}(\mathbf{w}^{(t)})$


- d. Increment t : $t \leftarrow t + 1$

- Output: $\mathbf{w}^{(t)}$


Linear Models




Linear Models




Linear Models




Linear Models?



Linear Models?



Nonlinear Models




Feature Transforms


- Given D -dimensional inputs $\mathbf{x} = [x_1, \dots, x_D]$, first compute some transformation of our input, e.g.,

$$\phi([x_1, x_2]) = [z_1 = (x_1 - 0.5)^2, z_2 = (x_2 - 0.5)^2]$$


Nonlinear Models



Nonlinear Models



Nonlinear Models



Nonlinear Models

