

10-701: Introduction to Machine Learning

Lecture 17 – Reinforcement Learning: Q-learning & Deep RL

Hoda Heidari

* Slides adopted from F24 offering of 10701 by Henry Chai.

Recall: Markov Decision Process (MDP)

- Assume the following model for our data:
 1. Start in some initial state s_0
 2. For time step t :
 1. Agent observes state s_t
 2. Agent takes action $a_t = \pi(s_t)$ $\pi: S \rightarrow A$
 3. Agent receives reward $r_t \sim p(r | s_t, a_t)$
 4. Agent transitions to state $s_{t+1} \sim p(s' | s_t, a_t)$
 3. Total reward is
$$\sum_{t=0}^{\infty} \gamma^t r_t \quad c \leq \gamma < 1 \text{ discount}$$
- MDPs make the *Markov assumption*: the reward and next state only depend on the current state and action.

Recall: Value Function

- $\underline{V^\pi(s)} = \mathbb{E}[\text{discounted total reward of starting in state } s \text{ and executing policy } \pi \text{ forever}]$

$$= \mathbb{E}[R(s_0, \pi(s_0)) + \gamma R(s_1, \pi(s_1)) + \gamma^2 R(s_2, \pi(s_2)) + \dots | s_0 = s]$$

$$= R(s, \pi(s)) + \gamma \mathbb{E}[R(s_1, \pi(s_1)) + \gamma R(s_2, \pi(s_2)) + \dots | s_0 = s]$$

$$= R(s, \pi(s)) + \gamma \sum_{s_1 \in \mathcal{S}} p(s_1 | s, \pi(s))(R(s_1, \pi(s_1)) + \gamma \mathbb{E}[R(s_2, \pi(s_2)) + \dots | s_1])$$

Recursive
Formula

$$V^\pi(\underline{s}) = \underbrace{R(s, \pi(s))}_{\text{Reward}} + \gamma \sum_{s_1 \in \mathcal{S}} \underbrace{p(s_1 | s, \pi(s))}_{\text{Transition Probability}} V^\pi(\underline{s_1}) \quad \forall s \in \mathcal{S}$$

Bellman equations

Recall: Optimality

- Optimal value function:

Recursive formula $V^*(s) = \max_{a \in \mathcal{A}} R(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, a)V^*(s')$

immediate *future*

- System of $|\mathcal{S}|$ equations and $|\mathcal{S}|$ variables
- Optimal policy:

$$\pi^*(s) = \operatorname{argmax}_{a \in \mathcal{A}} R(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, a)V^*(s')$$

*Greedy policy w.r.t v^** *Immediate reward* *(Discounted) Future reward*

Fixed Point Iteration

$$F: \mathbb{R}^n \rightarrow \mathbb{R}^n \text{ if fixed point: } F(x) = x \quad F(x_1, x_2, \dots, x_n) = (f_1(\vec{x}), f_2(\vec{x}), \dots, f_n(\vec{x}))$$

- Iterative method for computing fixed points of a fn
- Given some equations and initial values

$$\begin{cases} x_1 = f_1(x_1, \dots, x_n) \\ \vdots \\ x_n = f_n(x_1, \dots, x_n) \end{cases} \quad x_1^{(0)}, \dots, x_n^{(0)}$$

- While not converged, do

next estimate

$$\begin{cases} x_1^{(t+1)} \leftarrow f_1 \left(\underbrace{x_1^{(t)}, \dots, x_n^{(t)}}_{\text{current estimate}} \right) \\ \vdots \\ x_n^{(t+1)} \leftarrow f_n \left(\underbrace{x_1^{(t)}, \dots, x_n^{(t)}}_{\text{current estimate}} \right) \end{cases}$$

Fixed Point Iteration




Diagram from: <https://alexgs.co.uk/mathematics/unit-3-pure/3.8-numerical-methods/fixed-point-iteration>

Optimal Values as Fixed Points

Say $S = \{s_1, \dots, s_n\}$

- Optimal value function:

$$V^*(s) = \max_{a \in \mathcal{A}} R(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, a) V^*(s')$$

$$v^*(s_1) = \max_{a \in \mathcal{A}} \quad \text{(Diagram)} \quad s' \in \mathcal{S} \quad (s_1, \dots, s_n)$$

$$v^*(s_n) = \max \quad -a \quad (s_1, \dots, s_n)$$


$$v^*(s_1) = \overbrace{\{f\}}^{\text{Set}}(v^*(s_1), v^*(s_2), \dots, v^*(s_n))$$

$$v^*(s_n) \leq \underbrace{P_n}_{F} (v^*(s_1), v^*(s_2), \dots, v^*(s_n))$$

Value Iteration

- Inputs: $R(s, a)$, $p(s' | s, a)$
- Initialize $\underbrace{V^{(0)}(s)}_{\text{new estimate}} = 0 \forall s \in \mathcal{S}$ (or randomly) and set $t = 0$
- While not converged, do:
 - For $s \in \mathcal{S}$
$$\underbrace{V^{(t+1)}(s)}_{\text{new estimate}} \leftarrow \max_{a \in \mathcal{A}} R(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, a) \underbrace{V^{(t)}(s')}_{\substack{\text{prev. estimates} \\ Q(s, a)}}$$
 - $t = t + 1$
 - For $s \in \mathcal{S}$
$$\underbrace{\pi^*(s)}_{\substack{\text{Greedy} \\ \text{policy}}} \leftarrow \operatorname{argmax}_{a \in \mathcal{A}} R(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, a) V^{(t)}(s')$$
- Return π^*

Synchronous Value Iteration



- Inputs: $R(s, a), p(s' | s, a)$
- Initialize $V^{(0)}(s) = 0 \forall s \in \mathcal{S}$ (or randomly) and set $t = 0$
- While not converged, do:
 - For $s \in \mathcal{S}$
 - For $a \in \mathcal{A}$
$$\underbrace{Q(s, a)}_{\substack{= R(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, a)V^{(t)}(s')}}$$
 - $\underbrace{V^{(t+1)}(s)}_{\substack{\leftarrow \max_{a \in \mathcal{A}} Q(s, a)}}$
 - $\underbrace{t = t + 1}_{\substack{}}$
 - For $s \in \mathcal{S}$
$$\pi^*(s) \leftarrow \operatorname{argmax}_{a \in \mathcal{A}} R(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, a)V^{(t)}(s')$$
 - Return π^*

Asynchronous Value Iteration



- Inputs: $R(s, a), p(s' | s, a)$
- Initialize $V^{(0)}(s) = 0 \forall s \in \mathcal{S}$ (or randomly)
- While not converged, do:

- For $s \in \mathcal{S}$
 - For $a \in \mathcal{A}$

$$Q(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, a)V(s')$$

$$\cdot V(s) \leftarrow \max_{a \in \mathcal{A}} Q(s, a)$$

- For $s \in \mathcal{S}$

$$\pi^*(s) \leftarrow \operatorname{argmax}_{a \in \mathcal{A}} R(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, a)V(s')$$

- Return π^*

Value Iteration Theory

- **Theorem 1:** Value function convergence
 V will converge to V^* if each state is “visited” infinitely often (Bertsekas, 1989)
- **Theorem 2:** Convergence criterion
if $\max_{s \in \mathcal{S}} |V^{(t+1)}(s) - V^{(t)}(s)| < \epsilon$,
then $\max_{s \in \mathcal{S}} |V^{(t+1)}(s) - V^*(s)| < \frac{2\epsilon\gamma}{1-\gamma}$ (Williams & Baird, 1993)
- **Theorem 3:** Policy convergence
The “greedy” policy, $\pi(s) = \operatorname{argmax}_{a \in \mathcal{A}} Q(s, a)$, converges to the optimal π^* in a finite number of iterations, often before the value function has converged! (Bertsekas, 1987)

Bellman Optimality Characterization

- A policy π is optimal if and only if it is greedy (optimal) w.r.t. its own value function V^π .
- Proof:
 - (\Rightarrow) If π is optimal, then it must be greedy w.r.t V^π . If π were not greedy at some state, there would exist an action with strictly higher expected return \Rightarrow we could improve the policy $\Rightarrow \pi$ was not optimal. Contradiction.
 - (\Leftarrow) If π is greedy w.r.t V^π , then π is optimal.
Greedy w.r.t its own value solves the Bellman *optimality* fixed point, which is known to have a unique solution. So $V^\pi = V^*$ and π is optimal.

Policy Iteration

- Inputs: $R(s, a), p(s' | s, a)$
- Initialize π randomly
- While not converged, do:
 - Solve the Bellman equations defined by policy π

$$\underbrace{V^\pi(s)}_{\text{Bellman equation}} = R(s, \pi(s)) + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, \pi(s)) V^\pi(s')$$

- Update π

Greedy policy w.r.t V^π

$$\pi(s) \leftarrow \operatorname{argmax}_{a \in \mathcal{A}} R(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, a) \underbrace{V^\pi(s')}_{\text{Bellman equation}}$$

- Return π

Policy Iteration Theory

- In policy iteration, the policy improves in each iteration.
- Given finite state and action spaces, there are finitely many possible policies
 - Thus, the number of iterations needed to converge is bounded!
- Value iteration takes $O(|\mathcal{S}|^2|\mathcal{A}|)$ time / iteration
- Policy iteration takes $O(\underbrace{|\mathcal{S}|^2|\mathcal{A}|}_{\text{ }} + \underbrace{|\mathcal{S}|^3}_{\text{ }})$ time / iteration
 - However, empirically policy iteration requires fewer iterations to converge

Key Takeaways

- In reinforcement learning, we assume our data comes from a Markov decision process
- The goal is to compute an optimal policy or function that maps states to actions
- Value function can be defined in terms of values of all other states; this is called the Bellman equations
- If the reward and transition functions are known, we can solve for the optimal policy (and value function) using value or policy iteration
 - Both algorithms are instances of fixed point iteration and are guaranteed to converge (under some assumptions)

Two big Q's

1. What can we do if the reward and/or transition functions/distributions are unknown?
 s, a known
 $r(s, a)$ & $p(\cdot | s, a)$ unknown
2. How can we handle infinite (or just very large) state/action spaces?

Value Iteration

- Inputs: $R(s, a)$, $p(s' | s, a)$, γ
- Initialize $V^{(0)}(s) = 0 \forall s \in \mathcal{S}$ (or randomly) and set $t = 0$
- While not converged, do:
 - For $s \in \mathcal{S}$
 - For $a \in \mathcal{A}$
$$Q(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, a)V(s')$$
 - $V(s) \leftarrow \max_{a \in \mathcal{A}} Q(s, a)$
 - For $s \in \mathcal{S}$
$$\pi^*(s) \leftarrow \operatorname{argmax}_{a \in \mathcal{A}} R(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, a)V(s')$$
 - Return π^*

$Q^*(s, a)$ w/ deterministic rewards

Recursive
Formula

- $\underbrace{Q^*(s, a)}_{\text{total discounted reward}}$ = $\mathbb{E}[\text{total discounted reward of taking action } \mathfrak{a} \text{ in state } s, \text{ assuming all future actions are optimal}]$

$$= R(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, a) \underbrace{V^*(s')}_{\text{future value}}$$

$$\rightarrow V^*(s') = \underbrace{\max_{a' \in \mathcal{A}}}_{\text{maximize}} \underbrace{Q^*(s', a')}_{\text{future Q-value}}$$

$$Q^*(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, a) \left[\underbrace{\max_{a' \in \mathcal{A}}}_{\text{maximize}} \underbrace{Q^*(s', a')}_{\text{future Q-value}} \right]$$

$$\pi^*(s) = \operatorname{argmax}_{a \in \mathcal{A}} Q^*(s, a)$$

- Insight: if we know Q^* , we can compute an optimal policy π^* !

Q-learning: Overview

- Q-learning $\underset{\text{optimal}}{\checkmark} \quad Q^*(s, a) \quad \forall s \in S, a \in A$
 - Learn the Q-table without requiring a model of the world, that is rewards and transition probabilities (**model-free**).
 - Utilizes *Monte Carlo sampling*: it approximates the expectation simply by taking more and more samples over time.
 - Is **off-policy** because it estimates the value of the greedy optimal policy while following a potentially different policy.

$Q^*(s, a)$ w/ deterministic rewards

- $Q^*(s, a) = \mathbb{E}[\text{total discounted reward of taking action } a \text{ in state } s, \text{ assuming all future actions are optimal}]$
$$= R(s, a) + \gamma \sum_{s' \in \mathcal{S}} \underbrace{p(s' | s, a)}_{\text{red underline}} V^*(s')$$
$$V^*(s') = \max_{a' \in \mathcal{A}} Q^*(s', a')$$
$$Q^*(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} \underbrace{p(s' | s, a)}_{\text{red underline}} \left[\max_{a' \in \mathcal{A}} Q^*(s', a') \right]$$
$$\pi^*(s) = \operatorname{argmax}_{a \in \mathcal{A}} Q^*(s, a)$$
- Insight: if we know Q^* , we can compute an optimal policy π^* !

$Q^*(s, a)$ w/ deterministic rewards and transitions

- $Q^*(s, a) = \mathbb{E}[\text{total discounted reward of taking action } a \text{ in state } s, \text{ assuming all future actions are optimal}]$
$$= \underbrace{R(s, a)}_{\substack{\text{deterministic transition}}} + \gamma \underbrace{V^*(\delta(s, a))}_{s'}$$
- $\underbrace{V^*(\delta(s, a))}_{\substack{}} = \max_{a' \in \mathcal{A}} \underbrace{Q^*(\delta(s, a), a')}_{s'}$
- $$Q^*(s, a) = \underbrace{R(s, a)}_{\substack{}} + \gamma \max_{a' \in \mathcal{A}} \underbrace{Q^*(\delta(s, a), a')}_{\substack{}}$$
- $$\pi^*(s) = \operatorname{argmax}_{a \in \mathcal{A}} Q^*(s, a)$$
- Insight: if we know Q^* , we can compute an optimal policy π^* !

Learning $Q^*(s, a)$ w/ deterministic rewards and transitions

Algorithm 1: Online learning (table form)

- Inputs: discount factor γ , an initial state s
- Initialize $Q(s, a) = 0 \forall s \in \mathcal{S}, a \in \mathcal{A}$ (Q is a $|\mathcal{S}| \times |\mathcal{A}|$ array)
- While TRUE, do
 - Take a random action a
 - Receive reward $r = R(s, a)$
 - Update the state: $s \leftarrow s'$ where $s' = \underline{\delta}(s, a)$
 - Update $Q(s, a)$:

$$Q(\underline{s}, \underline{a}) \leftarrow r + \gamma \max_{a'} Q(\underline{s'}, \underline{a'})$$

Learning $Q^*(s, a)$ w/ deterministic rewards and transitions

Algorithm 2: ϵ -greedy online learning (table form)

- Inputs: discount factor γ , an initial state s ,
greediness parameter $\epsilon \in [0, 1]$ *exploration vs exploitation*

- Initialize $Q(s, a) = 0 \forall s \in \mathcal{S}, a \in \mathcal{A}$ (Q is a $|\mathcal{S}| \times |\mathcal{A}|$ array)

→ While TRUE, do

- With probability $\underline{1 - \epsilon}$, take the greedy action

$$a = \underset{a' \in \mathcal{A}}{\operatorname{argmax}} Q(s, a')$$

Otherwise, with probability ϵ , take a random action a

- Receive reward $r = R(s, a)$

- Update the state: $s \leftarrow s'$ where $s' = \delta(s, a)$

- Update $Q(s, a)$:

$$Q(s, a) \leftarrow r + \gamma \max_{a'} Q(s', a')$$

Learning $Q^*(s, a)$ w/ deterministic rewards

Algorithm 3: ϵ -greedy online learning (table form)

- Inputs: discount factor γ , an initial state s ,
greediness parameter $\epsilon \in [0, 1]$,
learning rate $\alpha \in [0, 1]$ (“trust parameter”)
- Initialize $Q(s, a) = 0 \forall s \in \mathcal{S}, a \in \mathcal{A}$ (Q is a $|\mathcal{S}| \times |\mathcal{A}|$ array)
- While TRUE, do

- With probability $1 - \epsilon$, take the greedy action

$$a = \underset{a' \in \mathcal{A}}{\operatorname{argmax}} Q(s, a')$$

Otherwise, with probability ϵ , take a random action a

- Receive reward $r = R(s, a)$
 - Update the state: $s \leftarrow s'$ where $s' \sim p(s' | s, a)$
 - Update $Q(s, a)$:

$$Q(s, a) \leftarrow (1 - \alpha) Q(s, a) + \alpha \left(r + \gamma \max_{a'} Q(s', a') \right)$$

Current value Update w/
deterministic transitions

Learning $Q^*(s, a)$ w/ deterministic rewards

Algorithm 3: ϵ -greedy online learning (table form)

- Inputs: discount factor γ , an initial state s ,
greediness parameter $\epsilon \in [0, 1]$,
learning rate $\alpha \in [0, 1]$ (“trust parameter”)
- Initialize $Q(s, a) = 0 \forall s \in \mathcal{S}, a \in \mathcal{A}$ (Q is a $|\mathcal{S}| \times |\mathcal{A}|$ array)
- While TRUE, do

- With probability ϵ , take the greedy action

$$a = \underset{a' \in \mathcal{A}}{\operatorname{argmax}} Q(s, a')$$

Otherwise, with probability $1 - \epsilon$, take a random action a



- Receive reward $r = R(s, a)$
 - Update the state: $s \leftarrow s'$ where $s' \sim p(s' | s, a)$ Temporal
 - Update $Q(s, a)$:

$$Q(s, a) \leftarrow \underbrace{Q(s, a)}_{\text{Current value}} + \underbrace{\alpha}_{\text{learning rate}} \left(r + \gamma \max_{a'} Q(s', a') \right) \underbrace{- Q(s, a)}_{\text{target}}$$


Temporal difference

difference

Learning $Q^*(s, a)$: Example




Learning $Q^*(s, a)$: Example




$Q(s, a)$	→	←	↑	↻
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0
5	0	0	0	0
6	0	0	0	0

Learning $Q^*(s, a)$: Example




Learning $Q^*(s, a)$: Example




$Q(s, a)$	→	←	↑	↻
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0
5	0	0	0	0
6	0	0	0	0


Learning $Q^*(s, a)$: Example



Learning $Q^*(s, a)$: Example



Learning $Q^*(s, a)$: Example



$$Q(3, \rightarrow) \leftarrow 0 + (0.9) \max_{a' \in \{\rightarrow, \leftarrow, \uparrow, \circlearrowright\}} Q(4, a') = 2.7$$

$Q(s, a)$	→	←	↑	↻
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	2.7	0	0	0
4	0	0	3	0
5	0	0	0	0
6	0	0	0	0

Learning $Q^*(s, a)$: Convergence

- For Algorithms 1 & 2 (deterministic transitions),
 Q converges to Q^* if
 - 1. Every valid state-action pair is visited infinitely often
 - Q-learning is exploration-insensitive: any visitation strategy that satisfies this property will work!
 - 2. $0 \leq \gamma < 1$
 - 3. $\exists \beta$ s.t. $|R(s, a)| < \beta \forall s \in \mathcal{S}, a \in \mathcal{A}$
 - 4. Initial Q values are finite (initialized to 0)


Learning $Q^*(s, a)$: Convergence

- For Algorithm 3 (temporal difference learning),
 Q converges to Q^* if
 - 1. Every valid state-action pair is visited infinitely often
 - Q-learning is exploration-insensitive: any visitation strategy that satisfies this property will work!
 - 2. $0 \leq \gamma < 1$
 - 3. $\exists \beta$ s.t. $|R(s, a)| < \beta \forall s \in \mathcal{S}, a \in \mathcal{A}$
 - 4. Initial Q values are finite
 - 5. Learning rate α_t follows some “schedule” s.t.
 $\sum_{t=0}^{\infty} \alpha_t = \infty$ and $\sum_{t=0}^{\infty} \alpha_t^2 < \infty$ e.g., $\alpha_t = \frac{1}{t+1}$

Two big Q's

1. What can we do if the reward and/or transition functions/distributions are unknown?
 - Use online learning to gather data and learn $Q^*(s, a)$
2. How can we handle infinite (or just very large) state/action spaces?

AlphaGo (Black) vs. Lee Sedol (White) Game 2 final position (AlphaGo wins)

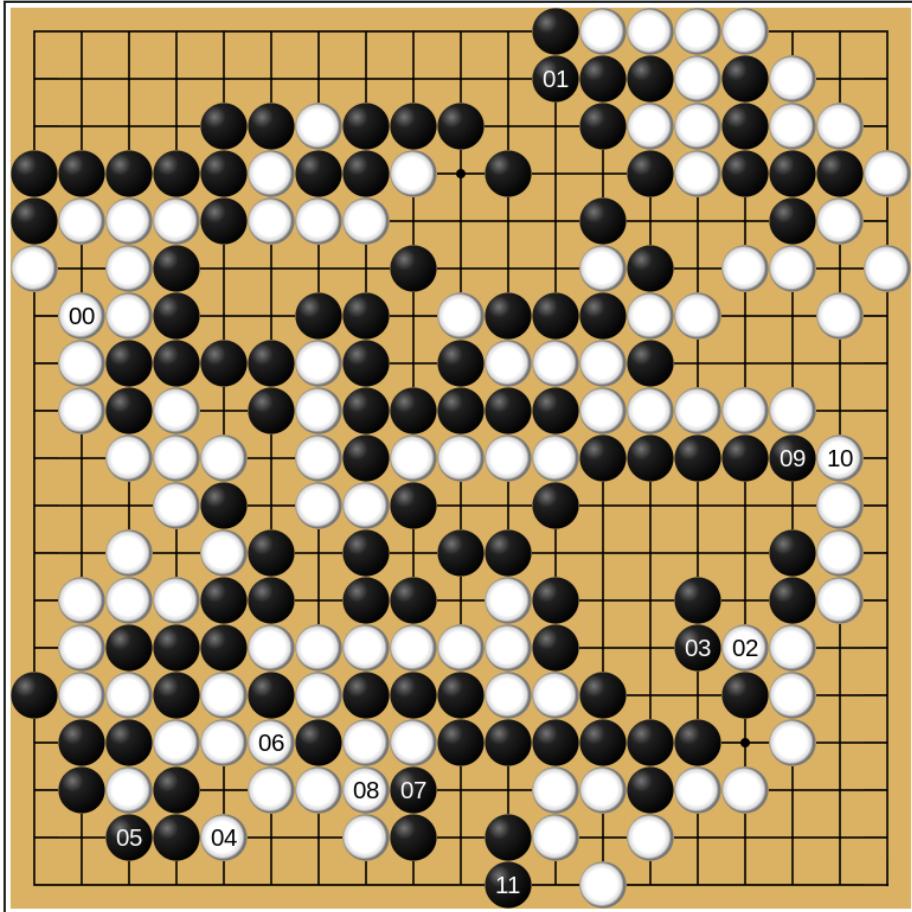


Source: https://en.wikipedia.org/wiki/AlphaGo_versus_Lee_Sedol

Playing Go

- 19-by-19 board
- Players alternate placing black and white stones
- The goal is claim more territory than the opponent
- How many legal Go board states are there?

AlphaGo (Black) vs. Lee Sedol (White) Game 2 final position (AlphaGo wins)



Henry Chai - 3/20/24

Source: https://en.wikipedia.org/wiki/AlphaGo_versus_Lee_Sedol

Source: https://en.wikipedia.org/wiki/Go_and_mathematics

Playing Go

- 19-by-19 board
- Players alternate placing black and white stones
- The goal is claim more territory than the opponent
- There are $\sim 10^{170}$ legal Go board states!

Two big Q's

1. What can we do if the reward and/or transition functions/distributions are unknown?
 - Use online learning to gather data and learn $Q^*(s, a)$
2. How can we handle infinite (or just very large) state/action spaces?
 - Throw a neural network at it!

Deep Q-learning


$Q^*(s, a) \quad \forall s \in S, a \in A$

parameter
↓

- Use a parametric function, $\underbrace{Q(s, a; \Theta)}$, to approximate $Q^*(s, a)$
 - Learn the parameters using *stochastic* gradient descent (SGD)
 - Training data $\underbrace{(s_t, a_t, r_t, s_{t+1})}$ gathered online by the agent/learning algorithm

Deep Q-learning: Model

- Represent states using some feature vector $s_t \in \mathbb{R}^M$
e.g. for Go, $s_t = [1, 0, -1, \dots, 1]^T$
- Define a *differentiable* function that approximates Q



Deep Q-learning: Loss Function

- “True” loss

$$\ell(\Theta) = \sum_{s \in \mathcal{S}} \sum_{a \in \mathcal{A}} \underbrace{(Q^*(s, a) - \underbrace{Q(s, a; \Theta)}_2)^2}_{\text{2. Don't know } Q^*}$$

1. \mathcal{S} too big to compute this sum

1. Use stochastic gradient descent: just consider one state-action pair in each iteration
2. Use temporal difference learning:

- Given current parameters $\Theta^{(t)}$ the temporal difference target is

$$Q^*(s, a) \approx r + \gamma \max_{a'} Q(s', a'; \Theta^{(t)}) := y$$

- Set the parameters in the next iteration $\Theta^{(t+1)}$ such that $Q(s, a; \Theta^{(t+1)}) \approx y$

$$\ell(\Theta^{(t)}, \Theta^{(t+1)}) = (y - Q(s, a; \Theta^{(t+1)}))^2$$

Deep Q-learning

Algorithm 4: Online learning (parametric form)

- Inputs: discount factor γ , an initial state s_0 , learning rate α
- Initialize parameters $\Theta^{(0)}$
- For $t = 0, 1, 2, \dots$
 - Gather training sample (s_t, a_t, r_t, s_{t+1})
 - Update $\Theta^{(t)}$ by taking a step opposite the gradient
$$\Theta^{(t+1)} \leftarrow \Theta^{(t)} - \alpha \nabla_{\Theta} \ell(\Theta^{(t)}, \Theta)$$
where
$$\nabla_{\Theta} \ell(\Theta^{(t)}, \Theta) = 2(y - Q(s, a; \Theta)) \nabla_{\Theta} Q(s, a; \Theta)$$

Deep Q-learning: Experience Replay

- SGD assumes i.i.d. training samples but in RL, samples are *highly* correlated
- Idea: keep a “replay memory” $\mathcal{D} = \{e_1, e_2, \dots, e_N\}$ of the N most recent experiences $e_t = (s_t, a_t, r_t, s_{t+1})$ (Lin, 1992)
 - Also keeps the agent from “forgetting” about recent experiences
- Alternate between:
 1. Sampling some e_i uniformly at random from \mathcal{D} and applying a Q-learning update (repeat T times)
 2. Adding a new experience to \mathcal{D}
- Can also sample experiences from \mathcal{D} according to some distribution that prioritizes experiences with high error (Schaul et al., 2016)

Key Takeaways

- We can use (deep) Q-learning when the reward/transition functions are unknown and/or when the state/action spaces are too large to be modelled directly
 - Also guaranteed to converge under certain assumptions
 - Experience replay can help address non-i.i.d. samples